



Free, Open Source, Offline Speech-to-Text

Geeks Club

 March 2026

Agenda

1.  **What is Handy?** — Overview & motivation
2.  **How It Works** — Workflow & key concepts
3.  **Architecture** — Tech stack deep dive
4.  **Speech Recognition Models** — Whisper & Parakeet
5.  **Installation & Usage** — Getting started
6.  **CLI & Platform Support** — Advanced usage
7.  **Conclusions & Discussion**

What is Handy?

A **cross-platform desktop application** that provides simple, privacy-focused speech transcription.


-  **Free** — Accessibility tooling belongs in everyone's hands
-  **Open Source** — MIT License, 82+ contributors
-  **Private** — Your voice stays on your computer, no cloud
-  **Simple** — One tool, one job: transcribe speech to text

"Your search for the right speech-to-text tool can end here — not because Handy is perfect, but because you can make it perfect for you."

Handy in Action

 [Add screenshot of Handy application here]

How It Works




Step by step:

1. **Press** a configurable keyboard shortcut (or use push-to-talk)
2. **Speak** your words while the shortcut is active
3. **Release** — Handy processes your speech using Whisper or Parakeet



Architecture Overview



Tech Stack Breakdown

Layer	Technology	Purpose
 Framework	Tauri v2	Cross-platform desktop app
 Frontend	React + TypeScript	Settings UI
 Styling	Tailwind CSS	UI styling
 Backend	Rust	System integration & ML inference
 STT (GPU)	whisper-rs	Whisper model inference
 STT (CPU)	transcription-rs	Parakeet model inference
 Audio	cpal	Cross-platform audio I/O
 VAD	vad-rs (Siler)	Voice Activity Detection
 Input	rdev	Global keyboard shortcuts
 Audio	rubato	Audio resampling

Language Split: Rust 49.1% | TypeScript 48.5% | Other 2.4%

Speech Recognition Models

Whisper Models (GPU-accelerated)

Model	Size	Best For
Small	487 MB	Quick transcription, lower resource usage
Medium (Q4)	492 MB	Balanced quality and speed
Turbo	1.6 GB	Fast, high quality
Large (Q5)	1.1 GB	Best accuracy

- Based on **OpenAI Whisper** via `whisper.cpp` / `ggml`
- GPU acceleration when available (NVIDIA, AMD, Intel, Apple Silicon)
- Quantized models (Q4, Q5) for smaller size with minimal quality loss

Parakeet V3 — CPU-Optimized Alternative

Why Parakeet?

-  **CPU-only** — no GPU required at all
-  **~5x real-time speed** on mid-range hardware (tested on i5)
-  **Automatic language detection** — no manual selection needed
-  **478 MB** compressed archive

Minimum Requirements

- Intel Skylake (6th gen) or equivalent AMD processors

Available Versions

Version	Size
Parakeet V2	473 MB

Installation

Download & Install

1. Download from [GitHub Releases](#) or [handy.computer](#)
2. **macOS**: Also available via Homebrew: `brew install --cask handy`
3. Grant system permissions (microphone, accessibility)
4. Configure keyboard shortcuts in Settings
5. Start transcribing!

Models

- Models are downloaded automatically on first use
- Manual download available for restricted networks
- Custom Whisper GGML models can be dropped into the `models` directory

Usage & Configuration

 [Add screenshot of keyboard shortcuts configuration here]

Key Features

- **Configurable shortcuts** — pick any key combo to start/stop recording
- **Push-to-talk mode** — hold to record, release to transcribe
- **Recording overlay** — visual indicator while recording
- **Debug mode** — `ctrl+Shift+D` (Windows/Linux) / `Cmd+Shift+D` (macOS)

CLI Parameters

Remote Control (send to running instance)

```
handy --toggle-transcription    # Toggle recording on/off
handy --toggle-post-process    # Toggle with post-processing
handy --cancel                  # Cancel current operation
```


Startup Flags

```
handy --start-hidden           # Start without main window
handy --no-tray                # Start without system tray icon
handy --debug                   # Enable verbose logging
```


Autostart Example

```
handy --start-hidden --no-tray # Silent background service
```

Platform Support



🤔 Handy vs. Alternatives



Handy's edge: GUI app + system integration + extensibility + active community




Key Takeaways

1. 🎙️ **Handy** is a free, open source, offline speech-to-text desktop app
2. 🦀 Built with **Tauri v2 + Rust** backend and **React + TypeScript** frontend
3. 🧠 Supports **Whisper** (GPU) and **Parakeet** (CPU-only) models
4. 🔒 **100% local** — no data leaves your machine
5. ⌨️ **Press → Speak → Release** — text appears in any app
6. 💻 Cross-platform: **macOS, Windows, Linux**
7. 🍴 Designed to be **"the most forkable"** speech-to-text app

Discussion

1. 🤔 Would you use an offline speech-to-text tool in your daily workflow?
2. 🦀 What do you think about **Tauri vs Electron** for desktop apps?
3. 🧠 **Whisper vs Parakeet** — when would you choose CPU-only over GPU?
4. 🔒 How important is **local/offline processing** for speech data?
5. 🛠️ Could we integrate something like Handy into our internal tooling?

Summary



Sources

-  [GitHub — cjpais/Handy](#)
-  [handy.computer](#) — Official website
-  [Discord Community](#)
-  [OpenAI Whisper](#) — Speech recognition model
-  [whisper.cpp](#) — C/C++ port of Whisper
-  [Tauri](#) — Desktop app framework
-  [Handy CLI](#) — Original Python CLI version

 **Thank You!**

Questions?

Handy — github.com/cjpais/Handy

 Join the community: [Discord](#)